

---

CONTENTS


EXERCISE A: NODE LEVEL STATISTICS . . . . .	<b>2</b>
EXERCISE B: RANDOM WALKS . . . . .	<b>3</b>
EXERCISE C: SHALLOW EMBEDDINGS . . . . .	<b>4</b>
EXERCISE D: PROGRAMMING EXERCISE . . . . .	<b>5</b>

---

---

**Exercise A** Node level statistics

Consider the following graph



The following code prints the adjacency matrix and computes its eigendecomposition:

```
>>> print(A)
[[0. 0. 1. 1. 0. 1. 0.]
 [0. 0. 0. 0. 1. 1. 1.]
 [1. 0. 0. 1. 0. 1. 0.]
 [1. 0. 1. 0. 0. 1. 0.]
 [0. 1. 0. 0. 0. 1. 1.]
 [1. 1. 1. 1. 1. 0. 1.]
 [0. 1. 0. 0. 1. 1. 0.]]

>>> lambda, E = np.linalg.eig(A)

>>> print(np.round(lambda, 3))
[ 3.646  2.   -1.646 -1.   -1.   -1.   -1.   ]

>>> print(np.round(E, 3))
[[-0.339 -0.408 -0.228 -0.816  0.004 -0.084  0.126]
 [-0.339  0.408 -0.228 -0.   -0.374  0.69  -0.255]
 [-0.339 -0.408 -0.228  0.408  0.511  0.114 -0.493]
 [-0.339 -0.408 -0.228  0.408 -0.515 -0.03  0.367]
 [-0.339  0.408 -0.228  0.   -0.176 -0.709 -0.377]
 [-0.558 -0.   0.83  -0.   0.   -0.   -0.   ]
 [-0.339  0.408 -0.228  0.   0.55  0.019  0.632]]
```

**Question A.1:** Determine the eigenvector centrality for each node in the graph.

**Question A.2:** Determine the clustering coefficient for each node in the graph.

**Question B.1:** Given a graph with adjacency matrix  $\mathbf{A}$  and a starting node chosen randomly according to a discrete distribution  $\mathbf{p}$ , what is the final node's probability distribution after taking a single step from the starting node along an edge chosen uniformly at random?

**Question B.2:** Given a graph with adjacency matrix  $\mathbf{A}$ , how many distinct paths of length  $t$  can we find starting from a specific node (say node 1)?

*Hint: We can represent the initial state as a vector that is one for the start node and zero elsewhere. Where can we end up after taking a single step, and how can this be computed using a matrix-vector product? How can this approach be generalized to  $t$  steps?*

In this exercise we will use the decoder  $\text{DEC}(\mathbf{z}_u, \mathbf{z}_v) = \sigma(\mathbf{z}_u^\top \mathbf{z}_v + b)$ , where  $\sigma(x) = \frac{1}{1+e^{-x}}$  denotes the sigmoid function.

**Question C.1:** As a warm-up, show that  $1 - \sigma(x) = \sigma(-x)$ .

Based on this, can you spot a typo in eq. 3.12 on page 27 in the book?

Let  $S_{u,v} \in \{0, 1\}$  denote a binary feature corresponding to an non-edge/edge between nodes  $u$  and  $v$ . Let  $P_{u,v} = P(S_{u,v} = 1 | \mathbf{z}_u, \mathbf{z}_v, b) = \sigma(\mathbf{z}_u^\top \mathbf{z}_v + b)$  denote the predicted probability that the edge is present, given the latent node embeddings  $\mathbf{z}_u, \mathbf{z}_v$  and bias  $b$ .

**Question C.2:** Write the cross entropy loss for the single observation  $S_{u,v}$  (in terms of  $P_{u,v}$  and  $S_{u,v}$ ).

**Question C.3:** Let us assume that the embeddings  $\mathbf{z}_u$  and  $\mathbf{z}_v$  are orthogonal, such that their dot product is zero, and let us further assume that the bias is zero,  $b = 0$ . What is the probability of an edge between node  $u$  and  $v$ ?

---

## Exercise D Programming exercise

In this exercise you will work with a shallow node embedding implemented in the script `shallow_embedding.py`. The code loads a graph from a file: This graph is simulated from a shallow embedding model, so that we know the ground truth probability of each possible link. In this exercise we will fit a shallow embedding model to the data and see how well we can estimate the ground truth.

**Question D.1:** Examine and run the code for loading the graph data.

- Understand how the graph is represented as a matrix as well as in the form of a set of index pairs and target values.
- It can perhaps help to visualize the adjacency matrix.

**Question D.2:** Examine and run the implementation of the class `Shallow`.

- Understand how the node embeddings are implemented using `torch.nn.Embedding`. Look up the documentation if needed.
- Understand what the forward function computes. What exactly is the role of the variables `rx` and `tx`?

**Question D.3:** Examine and run the code to fit the model. In this version, the loss is computed on the entire graph (no train/validation split and no mini batching).

- Experiment with different number of `max_step`.
- Experiment with different embedding dimensions. How does the embedding dimension influence the training loss?

**Question D.4:** Modify the code to use a train/validation split.

- Make a random split of the data (each node pair) into a training set (e.g. 80%) and a validation set (e.g. 20%).
- Modify the code to train on only the training data.
- Write code to compute the loss of the trained model on the validation set.
- Experiment with different embedding dimensions. What is the optimal embedding dimension when computing the loss on the validation set?

**Question D.5:** Hand in your predictions:

- Using the train/validation procedure you have implemented (or any other updates, hacks and modifications) to optimize the model. Compute what you believe is the best possible predicted link probability.
- Using the provided code, save your predictions in a file, `link_probabilities.pt`, and hand it in on DTU Learn.

I will compute the ground truth loss on your predictions and lowest generalization loss will be honored as the class winner.